

Exercises High Performance Data Networking

Niels van Adrichem

Chenxing Ji
Belma Turkovic

Fernando Kuipers¹
Adrian Zapletal

Jorik Oostenbrink

¹Responsible lecturer.

©2023 by Fernando Kuipers. All rights reserved. No part of the material protected by this copyright notice may be reproduced, redistributed, or utilized in any form or by any means, electronic or mechanical, including photocopying, recording or by any information storage and retrieval system, without written permission from Fernando Kuipers.

Chapter 1

Lecture 1: Introduction to Mininet

1.1 Important Notice

If you run into problems while going through any of the exercises, first try reading through the exercise again and fixing the problem by yourself. If this does not help, you can ask for assistance at the Q&A sessions or via the discussion forum.

At the end of the first and second chapters, you will find a summary of useful links and commands.

Another note: we have experienced that when copying commands or code directly from the reader (this document) into terminal, the terminal sometimes does not recognize certain special characters (e.g., ' or -).

1.2 Environment Setup

As you probably do not have access to a network of OpenFlow or P4 switches, you will run your programs on emulated networks instead. For this purpose we will use Mininet¹. Mininet emulates switches using Open vSwitch, a popular and open-source virtual switch that is used in both hardware switches sold by vendors and software switches that can be installed on generic computer hardware. In this section, we will download and set up a Virtual Machine containing a pre-installed version of Mininet and all other software required for the course exercises. We use different tools and VMs for x86- and ARM-based systems. Inside the VM, there should be no differences. If you have a Mac with M1 or M2 CPU, follow the instructions for ARM, otherwise x86. If you want to install everything yourself and not use the provided VM, you can download: <https://surfdrive.surf.nl/files/index.php/s/CSDP224IIId0eFpK>

x86 Install the open-source VirtualBox hypervisor, found here: <https://www.virtualbox.org/wiki/Downloads>.

The x86 VM image can be found here:

<https://drive.google.com/file/d/1oKpnLdGdJ2ETfo29UBguhDnaMaMef-5N/view?usp=sharing>

The Mininet and VirtualBox websites and communities provide ample information on installation issues. *We will not provide installation support for these tools.*

After downloading the image, start VirtualBox and import the VM by executing the following steps:

1. Import the VM by opening **File -> Import Appliance...** from the menu and selecting the image
2. Start the VM. The default username / password combination is: *hpdn / mininet*

ARM (Mac with M* CPU) Download the open-source QEMU-based hypervisor UTM (<https://github.com/utmapp/UTM>) from this link: <https://github.com/utmapp/UTM/releases/latest/download/UTM.dmg>

The ARM VM image can be found here:

<https://drive.google.com/file/d/1deICH4h2LsXIXbhyQ95BRz1lEjdIImog/view?usp=sharing>

After downloading and unzipping the image, start UTM and import the VM by executing the following steps:

1. Import the VM by selecting + and then *Open ...*; in this menu, select the downloaded image
2. Start the VM. The default username / password combination is: *hpdn / mininet*

1.3 Useful information about Mininet

Open a terminal, and start Mininet by running:

```
sudo mn
```

¹<http://mininet.org/>

By default, Mininet will start a virtual network with 2 hosts, h1 and h2, connected via switch s1. Run `?` to view all possible commands. Feel free to play around with them, e.g. run `pingall` to confirm connectivity between all hosts.

Each host has a separate network namespace (but hosts share access to all other system resources, such as the filesystem). Within Mininet, you can run most Linux commands directly on any of the virtual hosts by prepending the command by its hostname. For example, you can ping h2 from h1 by running `h1 ping h2` (Mininet automatically replaces the second “h2” with h2’s IP address). You can cancel the ping by pressing Ctrl+C. Analogously, you can also start different programs, services, or scripts on any of the virtual hosts, simply by prepending the relevant commands by a hostname.

Tip: By appending `&` to a command the process will run in the background and the Mininet window isn’t blocked.

You can emulate different network topologies by adding the `--topo` option to the start command for Mininet. Mininet itself comes with the following built-in topologies:

- **minimal:** The default topology of 1 switch with 2 hosts. No further parameters apply.
- **single:** A star topology with a single switch and `h` hosts. This topology can be used by appending `--topo single,h` to the `mn` command, where `h` refers to the number of hosts.
- **reversed:** Equivalent to the single switch topology, except that hosts connect to the switch in reverse order (i.e. the highest host number gets the lowest switch port).
- **linear:** `h` switches connect in a line, and there is one host connected to each switch. To use this topology, append `--topo linear,h` to the `mn` command.
- **tree:** A binary tree topology of depth `d`. To use this topology, append `--topo tree,d` to the `mn` command.

By default, hosts are assigned randomly generated MAC addresses. By appending `--mac` to the `mn` command, you can make Mininet use more readable MAC and IP addresses. This can be very helpful when debugging your controller application.

You can exit Mininet by typing `exit` in the terminal.

If Mininet crashes or you terminated it without using the `exit` command (for example with Ctrl+C), you will need to clean the Mininet environment with the following command:

```
sudo mn -c
```

For more information on how to create a Mininet network, you can run the following command to view the manual of Mininet:

```
man mn
```

1.4 Creating custom topologies

One of the advantages of Mininet is that it enables you to create complex network topologies and run experiments on them without the need to physically create those networks. You can specify your own custom network topologies in Python using a combination of the `addHost`, `addSwitch`, and `addLink` methods of the class `Topo`. To emulate your topology within Mininet, simply use the `--custom` command to load your custom Python script and `--topo` to select the (custom) topology:

```
sudo mn --custom /path_to_your_topology/topo_file.py --topo topology_name
```

For example, the following Python script creates a topology of 2 connected switches, adding 1 host to each switch, while the last line maps the topology classes to topology names. These names can then be used with the `--topo` option. The script is provided in the VM in `/home/hpdn/HPDN_Exercises/week_1/custom_topo.py`.

```
from mininet.topo import Topo
```

```
class MyTopo( Topo ):
```

```
    "Simple_topology_example."
```

```
    def __init__( self ):
```

```
        "Create_custom_topo."
```

```
        # Initialize topology
```

```
        Topo.__init__( self )
```

```
        # Add hosts and switches
```

```

leftHost = self.addHost( 'h1' )
rightHost = self.addHost( 'h2' )
leftSwitch = self.addSwitch( 's1' )
rightSwitch = self.addSwitch( 's2' )

# Add links
self.addLink( leftHost, leftSwitch )
self.addLink( leftSwitch, rightSwitch )
self.addLink( rightSwitch, rightHost )

```

```
topos = { 'mytopo': MyTopo }
```

It is also possible to map multiple topology classes to names in the same Python file, e.g.

```
topos = { 'mytopo': MyTopo, 'othertopo' : OtherTopo }
```

To start a network with the custom topology in this script, use:

```
sudo mn --custom /home/hpdn/HPDN_Exercises/week_1/custom_topo.py --topo mytopo
```

1.4.1 Exercise 1. - Complete graph & Square Lattice




Figure 1.1: Complete graph ($n = 5$).

The following script creates a complete topology of n switches, each with a single host (as in Figure 1.1):

```
from mininet.topo import Topo
```

```
class Complete(Topo):
```

```

    #switches form a complete graph
    #One host connected to each switch

```

```

    #We set a default value for n,
    #so Mininet will not crash if someone creates this topology without specifying n

```

```

def __init__(self, n = 5):
    Topo.__init__(self)

```

```
    switches = []
```

```
    for i in range(1,n+1):
```

```
        switch = self.addSwitch('s' + str(i))
```

```
        host = self.addHost('h' + str(i))
```

```
        self.addLink(switch, host)
```

```
    for s in switches:
```

```

self.addLink(switch, s)

switches.append(switch)

topos = {'complete': Complete}

```

The script is included in the VM: `/home/hpdn/HPDN_Exercises/week_1/custom_complete_topo.py`

For example, you can start a complete topology of 7 switches as follows:

```
sudo mn --custom /home/hpdn/HPDN_Exercises/week_1/custom_complete_topo.py --topo complete,7
```

Any parameters you provide for a custom topology (in this case 7) are passed on to `__init__`.

Note: Ping will not work on both the complete graph, as well as the square lattice topology you will construct yourself. (You will study this further in exercise 5 (Sec [1.8.2](#)))!

EXERCISE

Create a custom script that can construct square lattice topologies of arbitrary size $w \times w$. In a square lattice topology, all switches are aligned on a square grid and connected to their (up to 4) neighbors. Only add hosts to the four corner switches. See Figure [1.2](#) for an example of a square lattice topology of size 3×3 .




Figure 1.2: Square lattice graph ($w = 3$).

1.5 Introducing link properties

Mininet also has the capability to emulate network link parameters, such as bandwidth, delay, jitter, and loss. For example, if you want to set the bandwidth of all links in the network to 40 Mbps and their delay to 15 ms, you can run a command like the following. The default bandwidth unit is Mbps, and the default delay unit is μ s:

```
sudo mn --link tc,bw=40,delay=15ms
```

It is possible to add these parameters in the custom Python files as well by specifying these options in the `addLink` function and using `TCLink` as shown below, and this allows us to set parameters of each link separately:

```

from mininet.link import TCLink
...
self.addLink(s1, s2, delay='5ms', bw=20, cls=TCLink)

```

or

```

from mininet.link import TCLink
...
self.addLink(s1, s2, delay=5000, bw=20, cls=TCLink)

```

Instead of importing `TCLink` and using `cls=TCLink`, you can also add the `--link tc` option when running mininet. `tc` stands for traffic control.

1.5.1 Exercise 2.1 - Add link properties

Create a simple loop-free topology using a topology script. Make sure there is a path from `h1` to `h2` in your topology. You can use a topology like the initial example in Section 1.4 where 2 hosts are connected via 2 switches (`custom_topo.py`). Modify your topology to set the delay of each link to a random value between 0 ms and 10 ms. In addition, set all link bandwidths to 10 Mbps.

Confirm your results using `iperf` and `ping`, for example:

```
h2 iperf -s &
h1 iperf -c h2
h1 ping h2
```

Keep in mind that the resulting bandwidth and delay values might not be entirely precise because Mininet itself can be a little imprecise and we are running it inside a VM, causing more loss of precision.

1.5.2 Exercise 2.2 - Congestion Control

In this exercise, we will use Mininet's capability to set link properties to experiment with congestion control. For this, you first need to create a bottleneck, i.e., a link that has lower capacity than other links. This bottleneck is where congestion will happen. If you use `custom_topo.py`, you can create a bottleneck between the two switches. Set the bandwidth of the link between the two switches to 10 Mbps. Set the bandwidth of the links between the hosts and the switches to 100 Mbps. Fix the delay on every link to 5 ms. Additionally, set the buffer size at the bottleneck link to 25 packets using the parameter `max_queue_size=25` in the `addLink()` function.

Start the network and set the congestion control algorithm on one host (in this case `h1`) to Reno using

```
h1 ip route change 10.0.0.0/8 dev h1-eth0 congctl reno
```

Run `iperf` to send a data stream from this host to another host. The bandwidth you see should be close to 10 Mbps, which is the bottleneck link speed.

Now change the buffer size to 5 packets and restart the network. As before, set the congestion control algorithm to Reno and run `iperf`. What can you observe?

Finally, change the congestion control algorithm to BBR using

```
h1 ip route change 10.0.0.0/8 dev h1-eth0 congctl bbr
```

Run `iperf` again. What happens? Can you explain why?

1.6 Automating Tasks

It is possible to automate certain tasks in Mininet, such as adding routes to the hosts, executing Python scripts, or tearing down links. To do so, simply create a file and put a single command on each line, just like in a shell script. These commands can then be executed in Mininet with `source file`, where `file` is the path to your automation script.

For example, the following script first adds multicast routes to hosts `h1`, `h2`, and `h3` and prints "routes setup". Next, it starts two multicast `iperf` servers on hosts `h1` and `h2` and connects them with a client on host `h3`. After `h3` has finished sending its multicast packets, the two servers are shut down.

```
h1 route add -net 224.0.0.0 netmask 240.0.0.0 dev h1-eth0
h2 route add -net 224.0.0.0 netmask 240.0.0.0 dev h2-eth0
h3 route add -net 224.0.0.0 netmask 240.0.0.0 dev h3-eth0
py "routes_setup"
h1 iperf -s -B 224.0.0.14 -u -i 1 &
h2 iperf -s -B 224.0.0.14 -u -i 1 &
h3 iperf -c 224.0.0.14 -u -t 5
h1 kill %iperf
h1 kill %iperf
h2 kill %iperf
h2 kill %iperf
```

This script exists in `HPDN_Exercises/week_1/example_script`. You can create a single Mininet topology with 3 switches and run the script with:

```
sudo mn --topo single,3
```

```
mininet> source /home/hpdn/HPDN_Exercises/week_1/example_script
```

Important note: We put `&` after commands to start the process in the background. This way, the next command will execute immediately and would not have to wait until the earlier one is finished. In the case of `iperf`, we could alternatively use the `-D` option to run it as a daemon.

1.7 Custom Mininet commands

It is possible to create custom Mininet commands. For example, we can add the command “sleep” to Mininet by creating the following Python file and loading it with `--custom`:

```
from mininet.cli import CLI

from time import sleep

def custom_sleep(self, time):
    "custom_sleep_function"
    sleep(int(time))

CLI.do_sleep = custom_sleep
```

By specifying “`CLI.do_foo = bar`”, you create the custom command *foo* that executes the function *bar*.

If we load this custom file we can use `sleep 10` to let the Mininet command line interface sleep for 10 seconds.

You can have multiple custom files. For example, if we have a topology file “`custom_topo.py`” and a commands file “`custom_command.py`”, we can load both of them with:

```
--custom custom_topo.py,custom_command.py
```

1.7.1 Exercise 3. - Creating a Mininet script

Start Mininet with the linear topology with 4 switches (using `--topo=linear,4`) and the “`custom_command.py`” file created earlier, which contains the custom sleep command. Then, create an automation script that does the following:

- start a ping between h1 and h4
- sleep for 2 seconds
- tear down the link between switches s2 and s3 (`link s2 s3 down`)
- sleep for 60 seconds
- bring the link back up
- sleep for 2 seconds
- kill the ping process (`h1 kill %ping`)

Execute the script. What happens? What changes if you change the middle sleep command to 6 seconds instead of 60? Can you still deduce that the link was down from the ping messages?

Note that you only see the full output of ping **after** the process is killed.

1.8 Mininet packet capture using Wireshark

Wireshark is a network protocol analyzer used to capture and store network packets. Monitoring packets can be very helpful in later exercises to help you debug your application. In this section, we will use Wireshark to monitor the traffic between switches.

1.8.1 Exercise 4. - Monitoring traffic

1. Start Mininet with the linear topology with 2 switches.
2. Start Wireshark as a background process on host 1 (`h1 wireshark &`). Mininet always starts processes as superusers. This causes the initial warning you will see when starting Wireshark, and you can safely ignore this warning.
3. On the left you can select any of h1’s interfaces to monitor. We are only interested in `h1-eth0`, so select this one and press start. You should not see any packets yet, as there is currently no traffic in your emulated network.
4. Generate some traffic by starting a ping from host 1 to host 2. You should be able to spot 2 types of traffic: ARP packets (so h1 and h2 can learn each others’ addresses) and ICMP packets.
5. Filter on the ARP protocol by writing `arp` in the filter at the top-left.

6. You can filter packets on a wide variety of properties. For example, we can filter out all packets with IP destination address 10.0.0.2 by typing in `ip.dst!=10.0.0.2`. You can test this along with other conditions like the following:

```
ip.dst!=10.0.0.2 and icmp
```

The interfaces of switches can be monitored in the same way. For example, you can start Wireshark on switch s1 with `s1 wireshark &`. The only difference to hosts is that Mininet switches do not run on their own separate network namespace, so the Wireshark client we started on switch s1 can also access the interfaces of all other switches.

You can capture the traffic of multiple interfaces by either selecting the interfaces and pressing start, or by starting multiple Wireshark processes at the same host/switch and starting multiple captures.

Hint: When capturing multiple interfaces at the same time, you can add an additional column to indicate the interface the traffic was captured on by following these instructions: <https://osqa-ask.wireshark.org/questions/30636/traces-from-multiple-interface>

1.8.2 Exercise 5. - Diagnosing Network Issues

Start a Mininet instance with your square lattice topology (change the default width and set $w = 2$). Try to ping host h4 from host h2. Use Wireshark to find out what goes wrong and explain why this problem occurs.

Hint 1: Learning Switches. Switches start with an empty forwarding table and incrementally build it through a learning process based on the processed packets. This process works as follows: the switch maintains a table of MAC addresses that have appeared as source addresses in packets and the interface the switch has received these packets on. Thus, the switch learns and knows how to reach these addresses if one of them later shows up as a destination address. If a switch does not have an entry for a particular destination address, it defaults to flooding, i.e., it forwards the packet out on every port except the one on which it received the packet.

Hint 2: In previous exercises, we saw some network topologies where ping worked and others where it fails. What is a structural difference between the networks where ping succeeds and the ones where ping fails?

Hint 3: In Wireshark, order packets by timestamp. If your computer slows down too much during this exercise, restart the Mininet network.

Can you think of a method to fix the problem?

1.9 Useful commands

`dump`: Dump information about all nodes

`help`: Display all options

`device_name command`: If the first typed string is a host, switch or controller name, the command is executed on that node, e.g. to display the IP address of a host h1 type `h1 ifconfig`

`link s1 s2 down/up`: Bring links down and up

`h1 ping h2`: Test connectivity between two hosts

`pingall`: Test connectivity by having all nodes ping each other

`h1 arp`: Display the ARP table of host h1

`h1 route`: Display the routing table of host h1

`h1 python -m SimpleHTTPServer 80 &`: Run a simple web server on host h1

`h1 kill %python`: Shut the started web server down

`h2 wget -O - h1`: Send a request from node h2 to the web server running on node h1

You have now finished the HPDN Mininet exercises. For more information on Mininet, you can look through the official walkthrough: <http://mininet.org/walkthrough/>

