

# Quality of Service

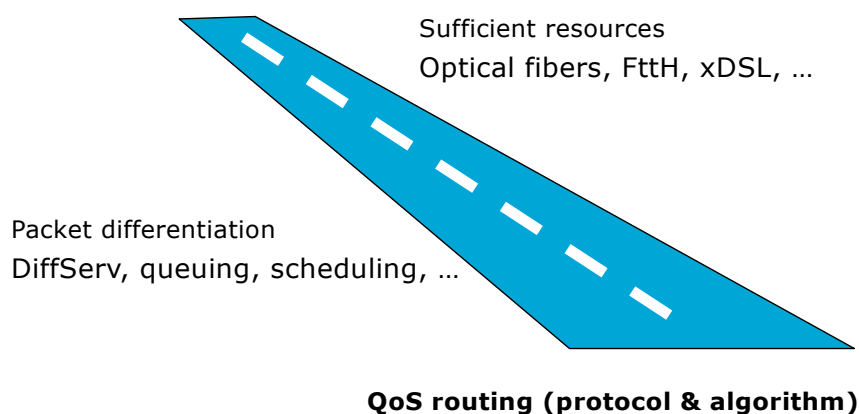
## Quality of Service (QoS)

- RFC 2386: "A set of service requirements to be met by the network while transporting a flow"
- Achieve certain, **specified levels of service**
- Parameters: Throughput, delay, jitter, packet loss, priority, ...
- Particularly important for **interactive** and **latency-sensitive** applications (telephony, gaming, video streaming, ...)

## Why do we need QoS?

- Without QoS: **Bufferbloat**
  - Router buffers fill up
  - Packets get delayed or dropped
- With QoS: certain packets get **treated differently**
- **Best-effort** service: default, no QoS mechanisms applied
- **Deterministic** service guarantees: network guarantees a parameter is kept below/above some threshold for certain application data (e.g., VoIP traffic delay < 5 ms)
- **Statistical** service guarantees: network guarantees that some percentage of traffic is kept below/above some threshold (e.g., for 80% of data delay < 20 ms)

## The road to QoS

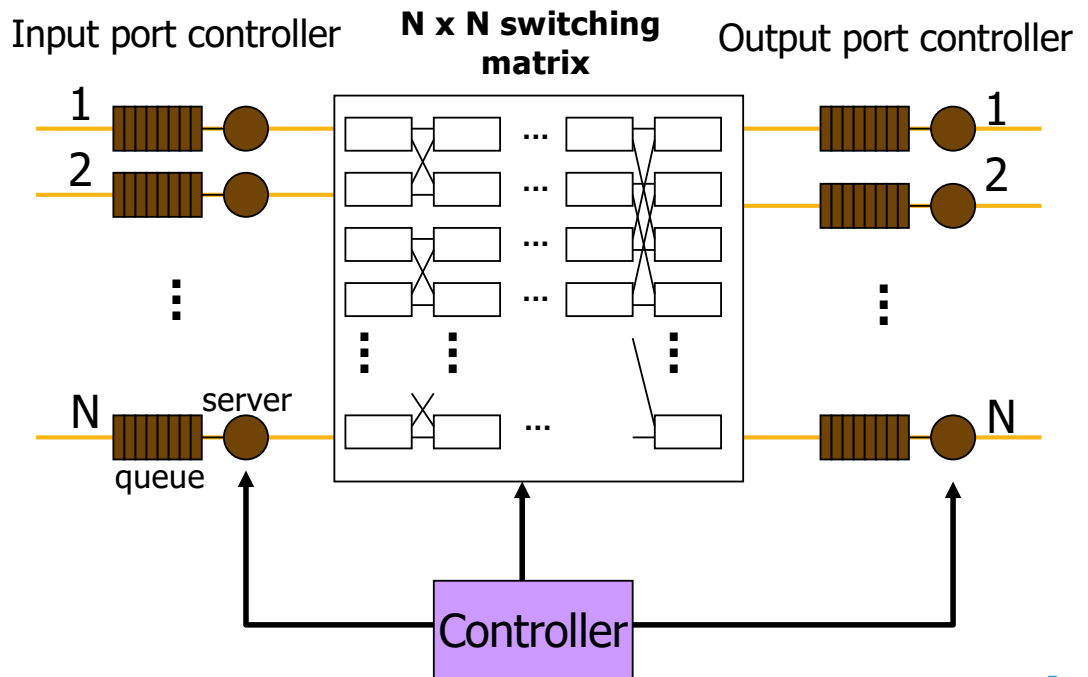


## What is required for achieving QoS?

- **Service Level Specifications**
  - Agreements between economic actors on the level of service that should be achieved
- **Classification** of flows into service classes
- **Input control**: traffic meters must ensure traffic conforms to expectations (e.g., limit data rate)
- **Scheduling** of packets at routers

# Scheduling

# Router



7

## Rule of thumb

- Short buffers, shorter delay, larger loss
- Large buffers, longer delay, smaller loss
- Loss and delay probabilities are increasing functions of the load


8

## Loss sensitive scheduling

- Assume 2 priority classes: high and low
- Packets in same class are served in FIFO order
- Buffer has K positions
- Head-of-the-Line, Partial Buffer Sharing, Push-out buffer, Random Early Detect

## Head of the Line (HoL)

The head of the line loss sensitive scheduling rule always serves high priorities in the buffer before low priorities:




Preemptive: placing back of low priority packet in server to queue  
(processing of a packet can get interrupted)

Non-preemptive: an arriving high priority packet has to wait until service of low priority packet has been terminated  
(processing of a packet cannot be interrupted)

Priority queuing: multiple priority classes

## Partial Buffer Sharing (PBS)

Above the threshold  $T$  only arriving high priority packets are allowed to enter and arriving low priority packets are discarded.




Below threshold: identical to FIFO and sequence order is preserved  
Above threshold: only a high priority regime until buffer is full

11

11

## Push-Out Buffer (POB)

Only if the buffer is full: an arriving high priority packet is allowed to push-out a previously entered low priority packet




LIFO POB: the last entered low priority packet is discarded  
FIFO POB: the first entered low priority packet is discarded  
R POB: a randomly chosen low priority packet is pushed out

12

12


# QoS guarantees

## Input Control




$V_{in}$  : Received rate  
 $V_{out}$  : Admitted rate  
 $V_c$  : Contracted rate

## Token Bucket ( $b, r$ )



## Burstiness Constraint: ( $\sigma, \lambda$ )





## Upper Bound Method

- Each customer arrival process satisfies a certain burstiness constraint
- The service time of each packet is deterministic (= proportional to packet length)
- Each non-empty QoS class  $k$  has a minimal service rate
  - Fair: no class can prevent another class from getting served

## Packets in the queue

- Burstiness constraint for class  $k$
- Service times are deterministic
- Fair scheduling policy
- Then queue count  $N_k(t)$  is bounded:

$$N_k(t) = \max_u [L_k(u, t) - M_k(u, t) : u \leq t]$$

$$N_k(t) = \max_u [\sigma_k + \lambda_k(t - u) - \mu_k(t - u) : u \leq t]$$

$$N_k(t) \leq \sigma_k : \mu_k \geq \lambda_k$$

$L$  is the number of inflow packets (upper-bounded by the burstiness constraint) and  $M$  is derived from the minimal service rate.

## Upper bound on delay & loss

- Assume single queue/multiplexer, H QoS classes.
- Zero packet loss if K is larger than sum of burstiness constraints  $\sigma_k$
- Delays are bounded by  $\sigma_k/\lambda_k$
- The deterministic approach shows that packet networks are able to guarantee loss and delay constraints.

## Reserving resources

## Resource reSerVation Protocol (RSVP)

- IETF's first signaling protocol (RFC-2205) based on multicast
- RSVP fundamental message types: *Path*, *Resv*
- *Path*:
  - previous hop,
  - sender template (describes the format of data packets that the sender will originate),
  - Tspec (specifies traffic characteristics of the sender's data),
  - Adspec (specifies the e2e QoS requirements)
- RSVP fundamental teardown messages: *PathTear*, *ResvTear*

## Resource ReSerVation Protocol

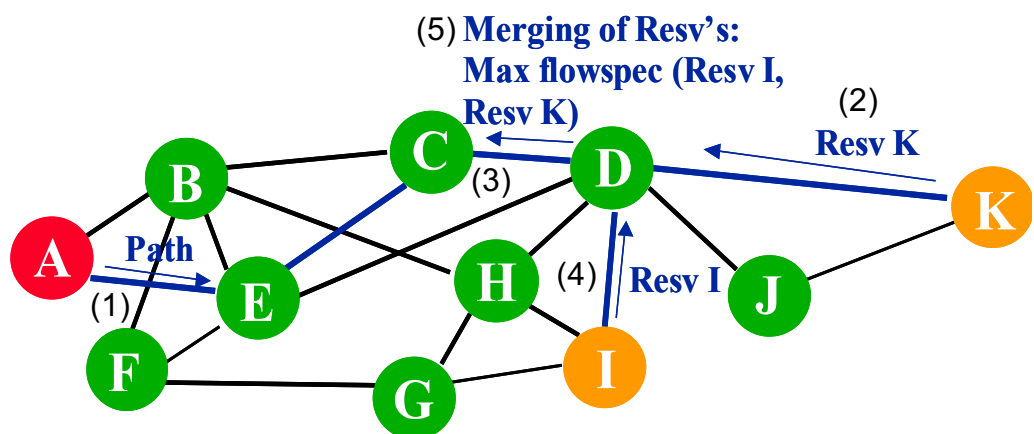
- RSVP is a *signaling* (and not routing) protocol
- Reservation for unicast as well as many-to-many multicast
- RSVP operates on top of IP (v4 or v6)
- RSVP model is receiver-oriented: receiver initiates and maintains reservation
- Receiver sends reservation requests upstream and each node either accepts or rejects it:
  - Router checks local policy and admission control
  - State info is stored in scheduler and classifier

## Resource reSerVation Protocol

- RSVP depends on routing protocol
- *Path* and *Resv* messages are sent **periodically** to maintain the reservation state along a particular traffic path

QoS state = SOFT STATE

## RSVP



# QoS protocols

IntServ & DiffServ

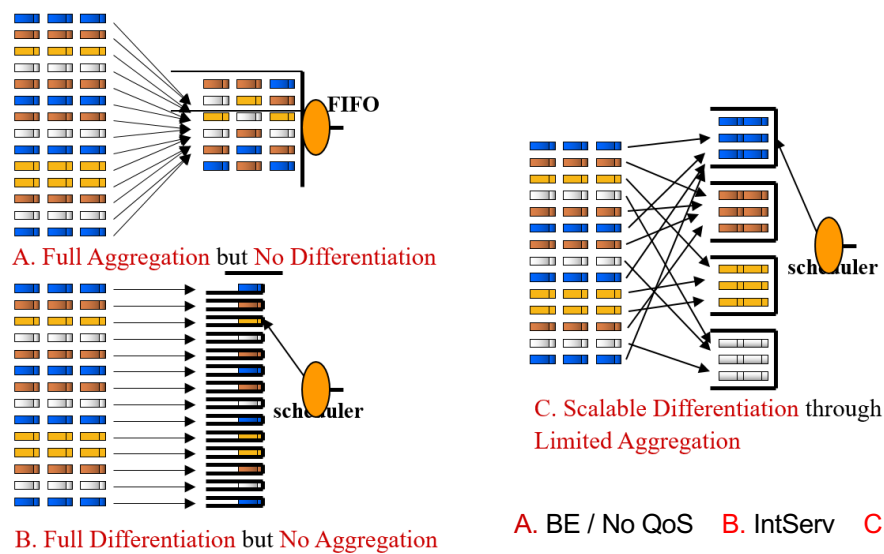
## Integrated Services (IntServ)

- Goal: provide guaranteed QoS
- Additional components:
  - packet classifiers (to identify flows that are to receive a certain level of service)
  - schedulers (to handle the service of different packet flows)
  - admission control (to determine whether a router has the necessary resources to accept a new flow)

## Integrated Services (IntServ)

- In combination with RSVP via (Tspec, Rspec)
- Traffic Classes:
  - Best-effort
  - controlled load (RFC2211): probabilistic guarantee
  - guaranteed service (RFC2212): on maximum queuing delay (not average, nor jitter)
- Disadvantage: complex and not scalable!
  - requires routers to keep state for each flow

## Differentiation / Aggregation



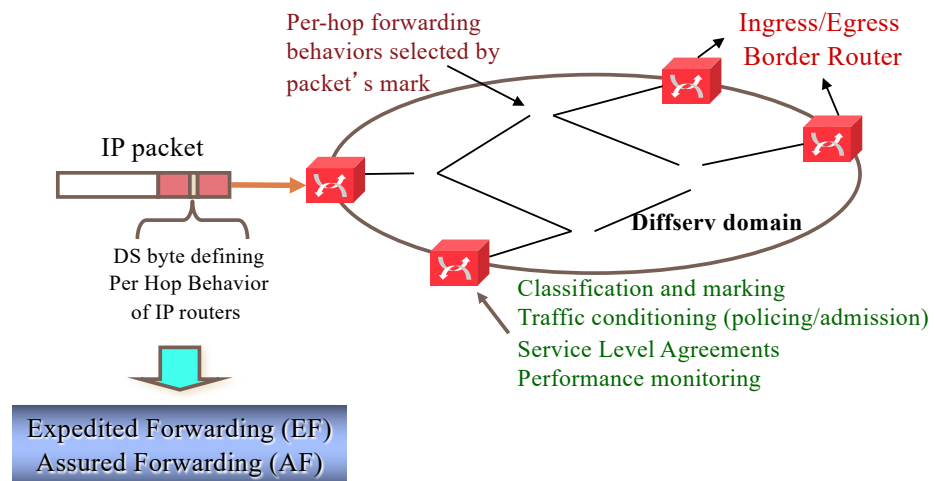
## Differentiated Services

- **Problem:** give *scalable* “better” service to some traffic (at the expense of worse service to the rest)
- IETF Differentiated Services (Diffserv) standardizes the “per hop behavior (PHB)” and not the service
  - service consists of “packet forwarding” + “rules”
- Types (besides Best Effort (BE)) of PHB
  - Expedited Forwarding (EF): virtual leased line where users want *absolute* BW independent of other traffic
    - implementation: priority queuing, strict policing
  - Assured Forwarding (AF): ‘better’ best effort with better control of *relative* BW
    - 4 AF classes, each with 3 drop preferences (green, yellow, red)
    - implementation: drop preference, weighted round-robin, WFQ
  - Other types exist, not covered here

## Basic difference between AF and EF

- Services built with EF PHB rely heavily on strict traffic conditioning at the network's boundary so that **excess traffic** and **congestion** within an EF Behavior Aggregate (BA) **is kept out of the network**
- Services built with AF PHB let **excess traffic** enter the network and use **active buffer management** at each DS network node **to handle congestion** within each AF class
- **Active** buffer management like **Random Early Drop (RED)** is recommended


## Diffserv Architecture



31

31

## Diffserv Code Points



The semantics of the IPv4 TOS byte (as well as IPv6 Traffic Class octet) are redefined

- **DSCP**: *Differentiated Services Code Point* index to identify/select the particular Per- Hop Behavior (PHB) an IP datagram is aiming to receive at a given network node
- **ECN**: *Not used for DiffServ but for Explicit Congestion Notification*
- **Unrecognized DSCP** should not be changed and should be forwarded transparently and **treated as Best Effort** traffic.
- **NCT** : Network Control Traffic

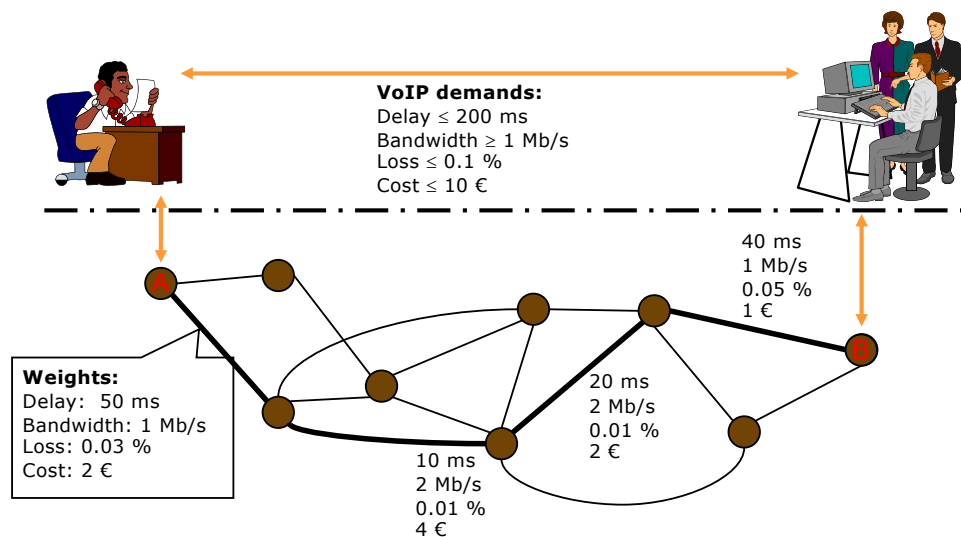
32

32



# QoS algorithm

## QoS routing problem (SAMCRA)




P. Van Mieghem and F.A. Kuipers, *Concepts of Exact Quality of Service Algorithms*, IEEE/ACM Transactions on Networking, vol. 12, no. 5, pp. 851-864, October 2004.

## QoS concepts in SAMCRA


- Non-linear length for exactness:

$$\max\left(\frac{w_1(P)}{L_1}, \frac{w_2(P)}{L_2}, \dots, \frac{w_m(P)}{L_m}\right)$$




Subsections of the shortest path are not always shortest paths


- We must compute  $k$ -shortest paths
- Reduce search space:
  - Non-dominance
  - Look-ahead



## Linear definition of path length



## Non-linear path length




$$\text{length}(P) = \left( \sum_{i=1, \dots, m} \left( \frac{w_i(P)}{L_i} \right)^q \right)^{1/q} \quad m = \text{number of constraints}$$

$$\lim_{q \rightarrow \infty} \left( \sum_{i=1, \dots, m} \left( \frac{w_i(P)}{L_i} \right)^q \right)^{1/q} = \max \left( \frac{w_1(P)}{L_1}, \frac{w_2(P)}{L_2}, \dots, \frac{w_m(P)}{L_m} \right)$$


## k-shortest paths

Example that subsections of the shortest path are not always shortest paths:



## Path Dominance

if  $w_i(P_1) \leq w_i(P_2)$  for all  $i$ , then  $P_1$  dominates  $P_2$




Reducing search space reduces complexity

## Look-Ahead

Compute for each of the  $m$  link weights the shortest paths tree rooted at B to any node  $n$  in network:

lower bounds  $b(n) = \{w_i(P_{B \rightarrow n}^*)\}$  for  $1 \leq i \leq m$

**Check:**  $w_i(P_{A \rightarrow n}) + b_i(n) \leq L_i$




## Meta-code of SAMCRA

SAMCRA( $G, m, A, B, L$ )

1. Initialize and find lower bounds
2. **While**( $Q$  not empty)
3.   Extract-min( $Q$ )  $\rightarrow u[i]$
4.   **if**  $u = B$ , **return** path
5.   **else for** each neighbor  $v$  of  $u$
6.     check for path dominance and length
7.     **if** length < maxlength and not dominated
8.       Insert the path in  $Q$
9.     **if**  $v = B$ , update maxlength

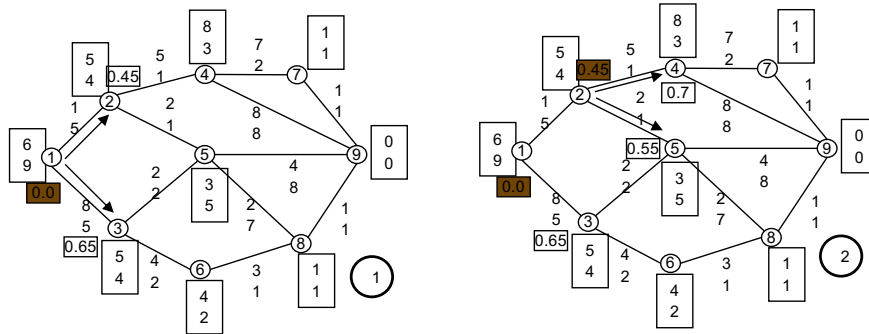
## Operation of SAMCRA



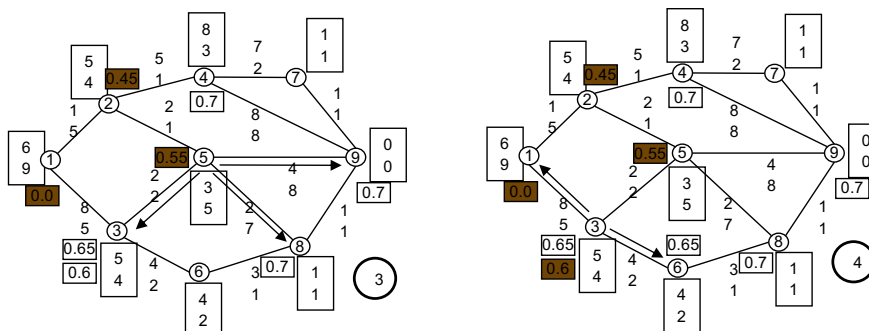
Source 1, destination 9, Constraints: (20,20)

Dijkstra lower bounds in rectangular boxes beside the node

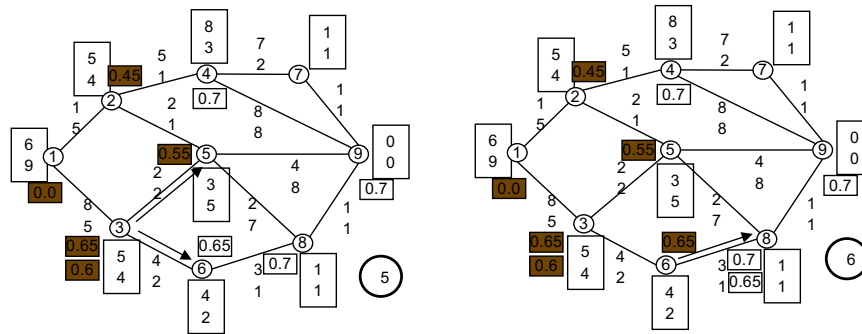
## Operation of SAMCRA



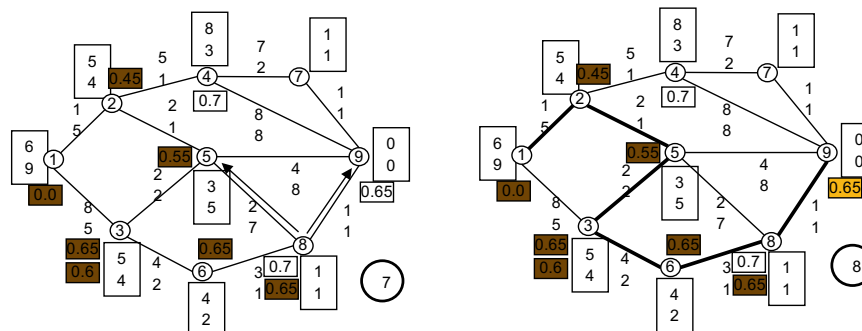
## Operation of SAMCRA



# Operation of SAMCRA



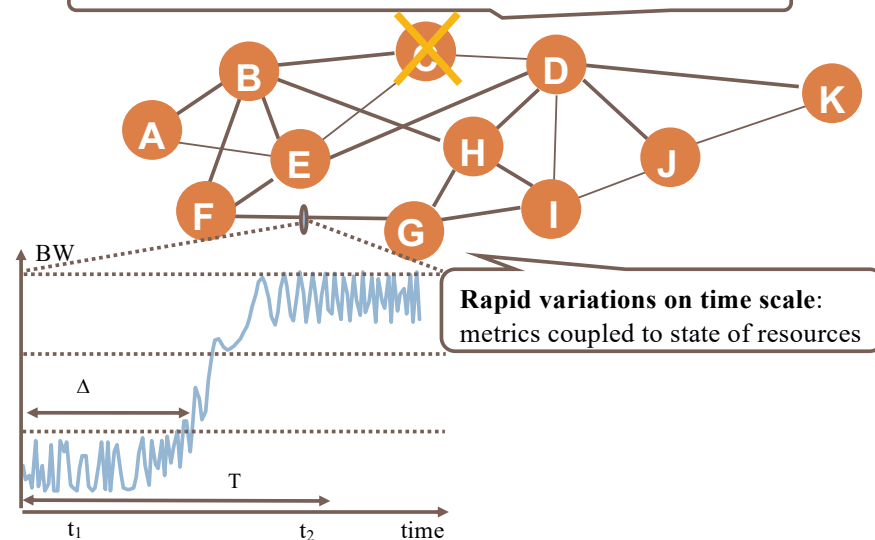
# Operation of SAMCRA



# QoS dynamics

## Topology Changes

Slow variations on time scale: failures, joins/leaves of nodes





## QoS dynamics

- Protocol:
  - How to distribute information on available resources?
  - When to distribute information on available resources?
- How to distribute:
  - Flooding -> may be costly if the frequency of updates is expected to be high
  - Tree-based protocols -> information is distributed over predetermined trees (less robust)

## When to distribute?


- Link-state update policies:
  - Periodic LSU policy
  - Trigger-based policy:
    - Class-based (equal or exponential classes)
    - Threshold
  - Additional techniques:
    - Hold-down timer
    - Moving average

## QoS with SDN

51

## Measuring throughput

- Per-flow counters
  - Packet counter
  - Byte counter
  - Flow duration



52


## Measuring throughput

- Traffic Switch 1 -> Switch 2

- $$Avg(Th) = \frac{ByteCounter}{FlowDuration}$$

- $$Th = \frac{\Delta BC}{\Delta FD}$$


- $Th_{Sw1} \geq Th_{Sw2} \geq Th_{Sw3} \geq Th_{Sw4}$ 
  - Due to packet loss



53

53

## Measuring packet loss




$$Avg(PacketLoss) = PacketCounter_{Sw1} - PacketCounter_{Sw4}$$

54

54

## Measuring delay

- Insert time-stamped probe packets
- Compare time-stamps at arrival
- Deduct delay between controller and switches



## Bandwidth guarantees with SDN

- OpenFlow supports 2 QoS mechanisms:
  - Queues
  - Meters

## Queues

- (Zero or more) queues bound to specific output port
- Parameters:
  - Minimum rate
  - Maximum rate
  - Priority
- Provide bandwidth guarantees

## Queues

- Cannot be configured with OpenFlow
- OpenFlow can only add flow entries placing packets in existing queues
- Configure queues directly with specific switch configuration protocol (e.g. OF-Config or Open vSwitch Database)

## Meters


- Configured with OF protocol
- Attached to flow entries
  - 1 meter can be attached to multiple flow entries
  - Can only attach 1 meter per flow entry
- Not properly supported by Open vSwitch

## Meters


- Measures packet rate of attached flow entries (via token bucket)
- Meter bands
  - Rate
  - Type
- Applies band with highest rate below the measured rate (if such a band exists)
- Types:
  - drop: drops packets
  - dscp remark: change DSCP field (differentiated services field)

## Illustration Queue & Meter

- Queue
  - Min\_rate
  - Max\_rate
  - Priority



- Meter
  - Drop
  - DSCP\_remark



## Possible design

- Admission control by the controller
- Traffic policing through *Meter table*
- Traffic prioritization at the switch through *Queue*
- Hard reservation

## Admission control

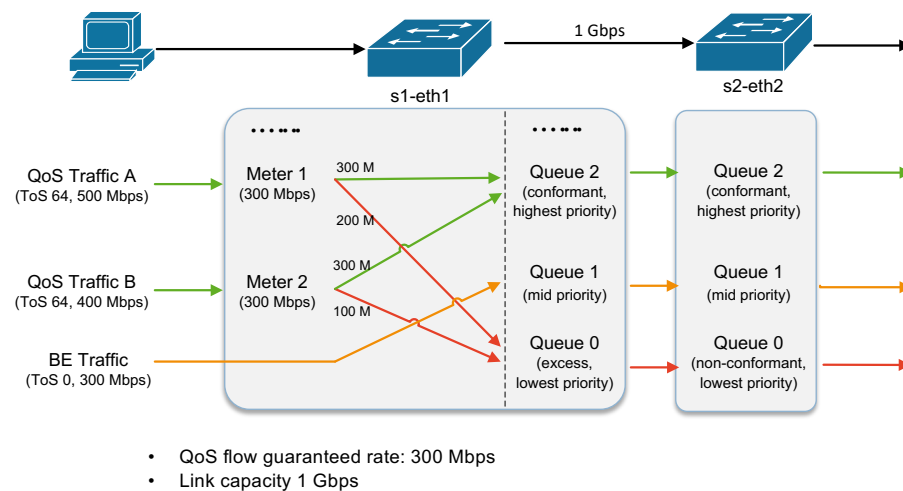
- Objective: to prevent QoS traffic from competing with each other
- QoS traffic requests to reserve bandwidth
- Controller rejects the QoS traffic if no path can accommodate the requested rate
- Best-effort traffic bypasses the admission control

## Traffic prioritization at the switch

- In each switch port, 3 queues:
  - Queue 2: Highest priority for conforming QoS traffic
  - Queue 1: Best-effort traffic
  - Queue 0: Excess QoS traffic



## Traffic aggregation



65

65

## Hard reservation

- QoS request to controller via OFPT\_PACKET\_IN with appropriate DSCP bits
- Resources freed via OFPT\_FLOW\_REMOVED
- Both are standard OpenFlow messages: no extra signaling overhead

66

66